

Seneca



CVI620/ DPS920

Introduction to Computer Vision

Image Segmentation and Deep Learning

Seneca College

Vida Movahedi

Overview

- Image vs. Object Segmentation
- Methods
 - Active Contour Methods
 - Perceptual Grouping or Contour Grouping Methods
 - Regional Methods
 - Region Growing
 - Grab-cut (graph-based)
 - Mean-shift
 - Deep Learning methods
- Introduction to Deep Learning




Image Segmentation

- Segmentation is a way to separate the image into regions with **homogenous behavior**
- Image segmentation is the task of finding groups of pixels that “go together” [1]
- Partition an image into regions, each corresponding to an **object** or **entity** [4]
- Each pixel is assigned a label such that pixels with the same label share certain properties, e.g. belonging to the same object (or background) [4]
- Many methods! Just mentioning a very few here!



- Various levels of detail
- Different when done by multiple human subjects



[Source: <https://www2.eecs.berkeley.edu/Research/Projects/CS/vision/bsds/>]

Object Segmentation


- Image segmentation is not the same as object segmentation.
- In salient **object segmentation**, the goal is to group pixels that belong to the **same object**;
- In other words, find the boundary between **foreground** (object) and **background**.

Segmentation: Two Perspectives! [4]




Figure 1.1: Salient object segmentation - (a) A sample image (source: [3]), (b) Segmentation as regional labeling, (c) corresponding object boundary.

How Can We Solve the Image Segmentation Problem?



<http://vladden.info/wp-content/uploads/FSO-1.jpg>




<https://i.ytimg.com/vi/7wCC2NaVLjs/maxresdefault.jpg>



https://vision.in.tum.de/_media/spezial/bib/ha_zirbas2014msc.jpg



https://vision.in.tum.de/_media/spezial/bib/ni_euwenhuis-cremers-pami12_2.jpg



<http://slideplayer.com/slide/4592950/15/images/16/Image+segmentation.jpg>

Active Contour Methods

Active Contour Methods

- Also called **snakes**
 - The active contour model is a method to **fit** open or closed splines **to lines or edges** in an image [5].
 - It works by **minimizing an energy** that is in part defined by the image and part by the spline's shape: length and smoothness. The minimization is done implicitly in the **shape energy** and explicitly in the **image energy**. [6]
-
- [5] Snakes: Active contour models. Kass, M.; Witkin, A.; Terzopoulos, D. International Journal of Computer Vision 1 (4): 321 (1988). DOI:10.1007/BF00133570
 - [6] http://scikit-image.org/docs/dev/auto_examples/edges/plot_active_contours.html

Snakes and Energy

- A simple **elastic** snake is defined as a curve


$$f(s) = (x(s), y(s)), s \in [0,1]$$

- Goal: minimize the energy of the snake

$$E_{\text{snake}} = E_{\text{internal}} + E_{\text{external}} = E_{\text{internal}} + E_{\text{image}} + E_{\text{con}}$$

E_{internal} :	Controls the deformation of the snake. A smooth and continuous snake has a lower internal energy than a stretched snake or one with high curvature (the elasticity and smoothness)
E_{image} :	Controls the fitting of the contour to the image. A snake which passes through dark ridges, edges with strong gradient (high magnitude), or corners and line terminations has a lower image energy than one that does not pass through these (how well the contour contracts or expands to the edges)
E_{con} :	User constraints or prior model constraints. A snake that is close (or away) from user specified features has a lower constraint energy than one that does not

Click to select initial contour location. Double-click to confirm and proceed.



<https://www.youtube.com/watch?v=3-YQ9k47ktY>


Region Growing

Grab-Cut (Graph-Based)

Mean-Shift

Contour Grouping Methods

- Grouping edges or line approximations of edges into a closed boundary of an object [4]



Region Growing Methods [2]


- Flood Fill

- A seed point is selected from an image
- All *similar* neighboring pixels are labelled with the **same label** (belonging to the same region)
- **Similarity** is defined as being within a range of the **seed** or **neighbors**

Procedure

- A seed point is selected, then all **similar** neighboring points are colored with a uniform color
- Neighboring pixels need not be identical in color, but within a specified **range** (loDiff to upDiff) of either the current pixel or the original seed value

Region Growing Methods



<https://www.youtube.com/watch?v=SokLiR6PA0I>

OpenCV Flood Fill [2]

```
int cv::floodFill(  
    cv::InputOutputArray image,           // Input image, 1 or 3 channels  
    cv::Point seed,                       // Start point for flood  
    cv::Scalar newVal,                    // Value for painted pixels  
    cv::Rect* rect,                       // Output bounds painted domain  
    cv::Scalar lowDiff = cv::Scalar(),    // Maximum down color distance  
    cv::Scalar highDiff = cv::Scalar(),    // Maximum up color distance  
    int flags                             // Local/global, and mask-only  
);
```

Python:

```
cv.floodFill(image, mask, seedPoint, newVal[, loDiff[, upDiff[,  
flags]]) -> retval, image, mask, rect  
mask: to control where filling is done
```

Region Growing Methods [2, 5]

Watershed

- It is useful if we do not have the benefit of separate background mask
- An image is interpreted as a height field or landscape. Uniform areas are **valleys**, lines are **mountains**
- Fill the valley with water, until water from two valleys are about to merge
- Build a barrier to prevent merging
- Resulting barrier is the image segmentation

Procedure

- Takes the gradient of intensity to form valleys and mountains
- Start flooding the landscape at all local minima
- Wherever different evolving components meet, label the ridges (barriers)

OpenCV Watershed [2]

- An implementation of the watershed algorithm
- Allows the user to mark parts of an object or background
- It tells the algorithm to “group points like these together”




Figure 12-12. Watershed algorithm: after a user has marked objects that belong together (left), the algorithm merges the marked area into segments (right)

Watershed (cont.)

```
void cv::watershed(  
    cv::InputArray      image,           // Input 8-bit, three channels  
    cv::InputOutputArray markers        // 32-bit float, single channel  
);
```

Python: `cv.watershed(image, markers) -> markers`

- The second argument, markers, is a single-channel integer (cv::S32) image, having the same dimension, and 0 everywhere except at caller's marks
- For example, the orange might have been marked with a "1," the lemon with a "2," the lime with "3," the upper background with "4," and so on
- The function will set all 0 pixels of markers to one of the marked values, except the boundaries between regions, which may be set to -1.


Region Growing

Grab-Cut (Graph-Based)

Mean-Shift

Grab-Cut [2, 5]

- A graph method: capable of obtaining excellent segmentations, often with just a bounding rectangle around the object
- It uses user-labeled foreground and background regions to establish distribution histograms
- Unlabeled foreground or background should conform to similar distributions (smooth and connected)
- Assertions are combined into an energy functional: a low energy to solutions that conforms to these assertions and a high energy to solutions that violates them.
- Final result is by minimizing this energy function



OpenCV Grab-Cut

```
void cv::grabCut(  
    cv::InputArray      img,  
    cv::InputOutputArray mask,  
    cv::Rect            rect,  
    cv::InputOutputArray bgdModel,  
    cv::InputOutputArray fgdModel,  
    int                 iterCount,  
    int                 mode      = cv::GC_EVAL  
);
```

Python:

```
cv.grabCut( img, mask, rect, bgdModel, fgdModel,  
            iterCount[, mode] ) -> mask, bgdModel, fgdModel
```

Region Growing

Grab-Cut (Graph-Based)

Mean-Shift

OpenCV Mean-Shift

- Mean-Shift looks at the **spatial** distribution of color (x,y)
- OpenCV implementation color, replace value with **mean**
- Segmentation is done over a scale pyramid
- Output is a “posterized” image (fine texture is removed and gradient is color are flattened)

Python:

```
cv.pyrMeanShiftFiltering  
(src, sp, sr[, dst[,  
maxLevel[, termcrit]])  
    ) -> dst
```

```
void cv::pyrMeanShiftFiltering(  
    cv::InputArray  src,                // 8-bit, Nc=3 image  
    cv::OutputArray dst,                // 8-bit, Nc=3, same size as src  
    cv::double      sp,                 // Spatial window radius  
    cv::double      sr,                 // Color window radius  
    int              maxLevel = 1,      // Max pyramid level  
    cv::TermCriteria termcrit = TermCriteria(  
        cv::TermCriteria::MAX_ITER | cv::TermCriteria::EPS,  
        5,  
        1  
    )  
);
```

```
cv::pyrMeanShiftFiltering( src, dst, 20, 40, 2);
```



Figure 12-13. Mean-shift segmentation over scale using `cv::pyrMeanShiftFiltering()` with parameters `max_level=2`, `spatialRadius=20`, and `colorRadius=40`; similar areas now have similar values and so can be treated as super pixels (larger statistically similar areas), which can speed up subsequent processing significantly

New Methods: Deep Learning Semantic Image Segmentation – Microsoft Ignite



<https://youtu.be/FroRjEejA30>

New Methods: Deep Learning

Segment Anything – Meta AI CV Research

- *Segment Anything Model (SAM)* is a new AI model from Meta AI that can "cut out" any object, in any image, with a single click
- *The Project is* a new task, model, and dataset for image segmentation
- <https://segment-anything.com/>
- https://www.youtube.com/shorts/oYUcl_cqKcs



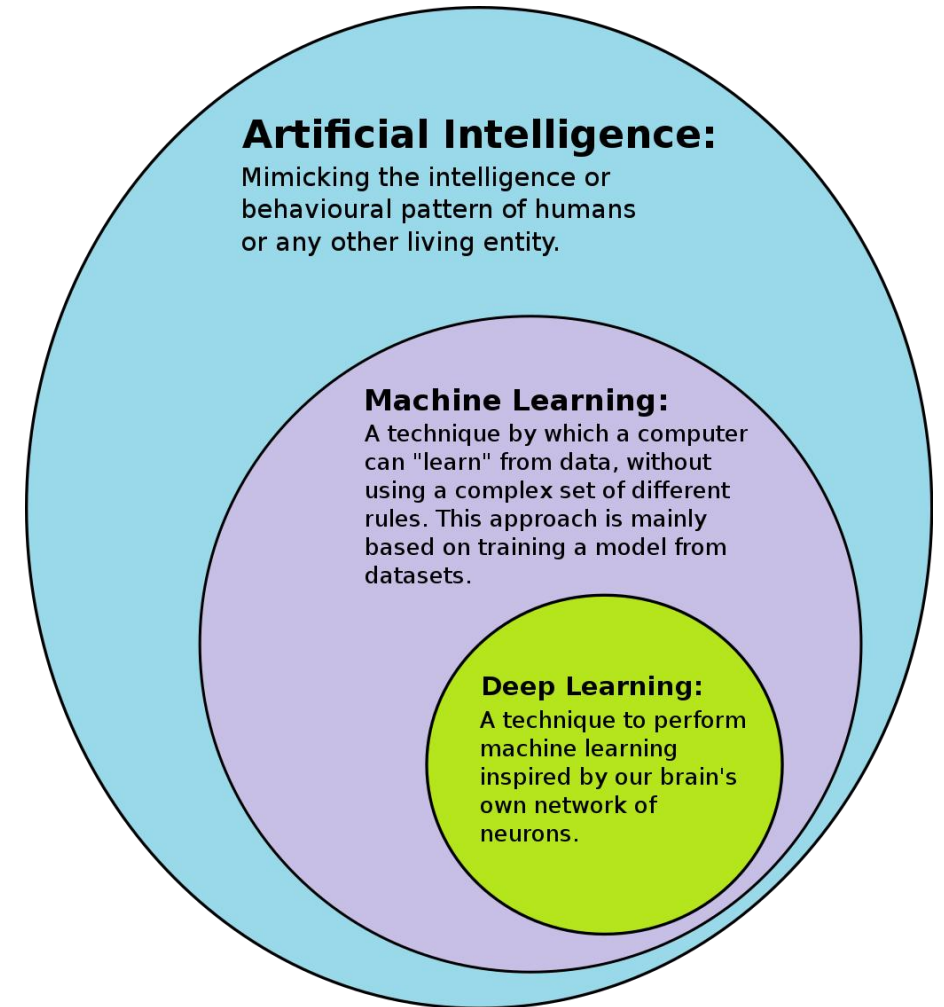
Quick Introduction to Deep Learning: DL-Based (Method 3)

Reminder

- Computer Vision problems can be solved using three approaches:
 - 1) Start from intuition and design a method, or
 - 2) Start from data and use Machine Learning, or
 - 3) Start from LOTS of data and use Deep Learning

Deep Learning


- A technique for implementing Machine Learning
- Allows for learning models
- Learning algorithm that is inspired by how biological brains work



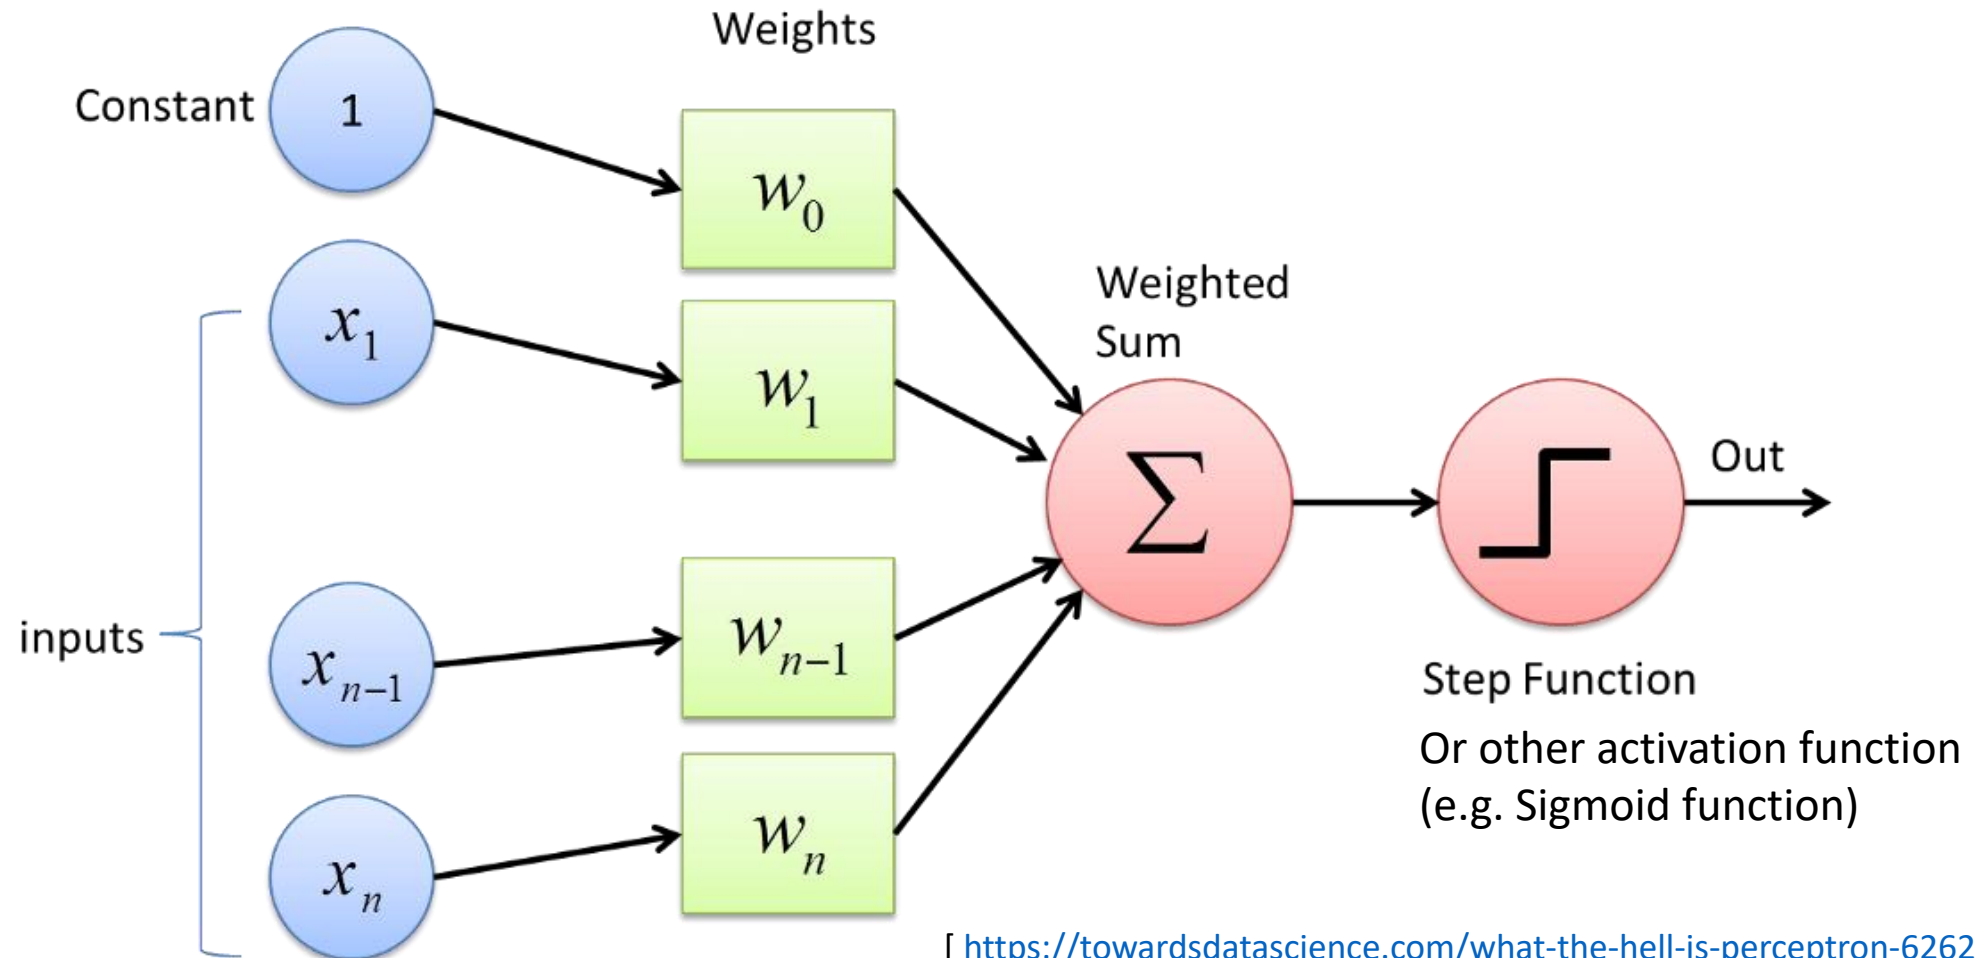
Neuron

Human brains have millions of neurons; they receive, process, and transmit electric and chemical signals

- Nucleus
- Dendrites: receives signals from other neurons
- Axon: a single output filament
- Synapse: acts as communicator




Artificial Neuron (Perceptron or Single Layer Neural Network)



[<https://towardsdatascience.com/what-the-hell-is-perceptron-626217814f53>]


Perceptron for Classification

- Same as logistic regression
- A Perceptron can be trained to perform **binary** classification, when classes are **linearly separable**
- By tuning the weights, we can train the perceptron



Biological Neural Networks

- Cerebral Cortex [Wikipedia]
 - Outer layer of neural tissue of the cerebrum in humans & mammals
 - Key role in attention, perception, awareness, thought, memory, language, and consciousness.
 - Mostly consists of the six-layered neocortex





https://www.frontiersin.org/files/Articles/285579/fnana-11-00102-HTML/image_m/fnana-11-00102-g002.jpg

Artificial Neural Networks

- A network of neurons, connected to each other in a certain way
- Multilayer Perceptron (MLP)
- Most popular structure is **feed-forward** and with **hidden** layers


[https://www.researchgate.net/figure/The-structure-of-SOM-neural-network_fig1_271637449]



Learning the Weights

- How do we figure out the weights in the network?
- Training or learning
- Just like a child learning from experience and feedback, a network can be trained
 - Formation and modification of synapses between neurons
 - Trial and error
- Learning algorithms

Neural Net Training- Backpropagation

- First, all weights initialized to small random values
- Loop
 - For all training samples
 - epoch 
 - Pass through the network, calculate prediction
 - Calculate the error between prediction and actual target value
 - Back-propagate the error
 - Update weights proportional to the error
- Until convergence to a minimum error

Issues for Learning in Deep Networks


- **Deep Networks**
 - Having many and many layers
- **Vanishing / Exploding Gradient**
 - When backpropagating through a deep network, the signal quickly fades to near zero values
 - Activation functions need large enough values to let the signal pass; therefore, low values are blocked.
 - The farther neuron layers are from the output, the higher the likelihood that they'll get locked out of updates.
 - Therefore, the network stops learning as a whole (or learns at incredibly slow pace).

Issues for Learning in Deep Networks (cont.)

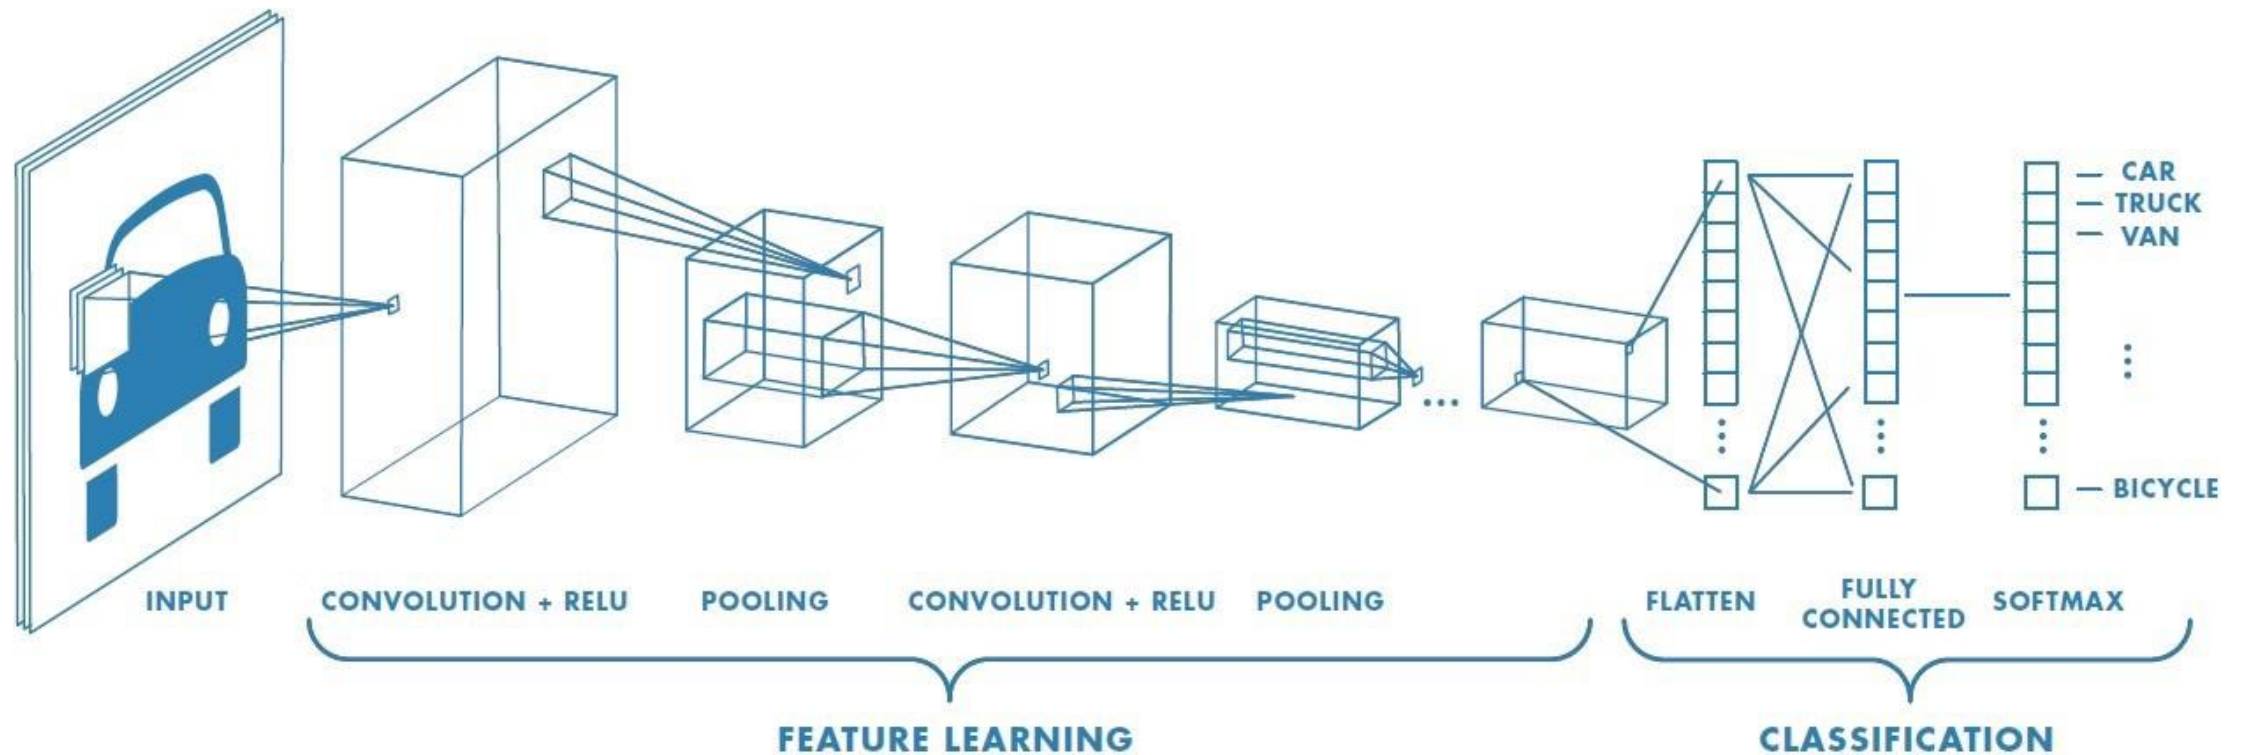
- New improvements in Neural Net training, overcoming previous shortcomings, that allows for training of **deep** networks.
 - Improvements in algorithm
 - Geoffrey Hinton, UofT (now Google)
 - ImageNet Classification with Deep Convolutional Neural Networks: Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton 2012

Deep Convolutional Neural Nets (CNN)


- Invariance
 - Having the same response regardless of size, distortion, or position in the image
- Based on “Convolution”
 - Applied similarly to all parts (tiles) of images



Example




<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>



Open-Source Frameworks

- Share the **trained** neural nets
 - NN Architecture
 - Learned weights
- Democratization of Deep Learning models




YOLO V8

[<https://yolov8.com/>]

Classification	Detection	Segmentation
		

MASK R-CNN


[https://github.com/matterport/Mask_RCNN]



GoogLeNet (Inception)


[<https://www.kaggle.com/keras/inceptionv3/home>]

Image



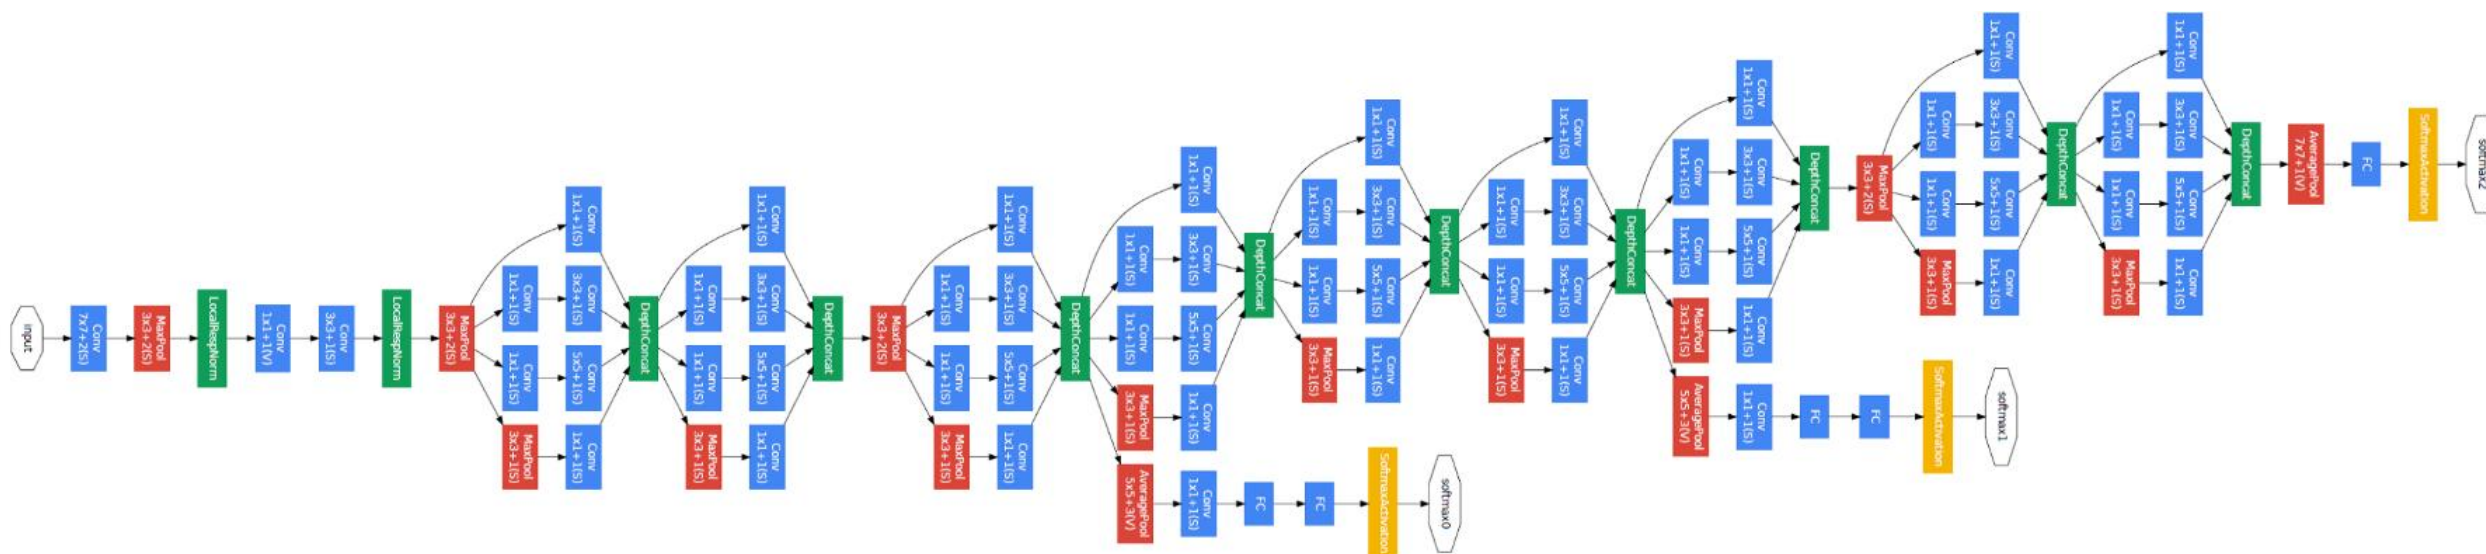
Inception Labels

1. Shih-Tzu: 79.418%
2. Lhasa: 2.359%
3. Affenpinscher: 1.567%
4. Pekinese: 0.954%
5. Tibetan Terrier: 0.661%



1. Assault Rifle: 87.250%
2. Rifle: 6.149%
3. Bulletproof Vest: 3.534%
4. Military Uniform: 1.273%
5. Stretcher: 0.060%

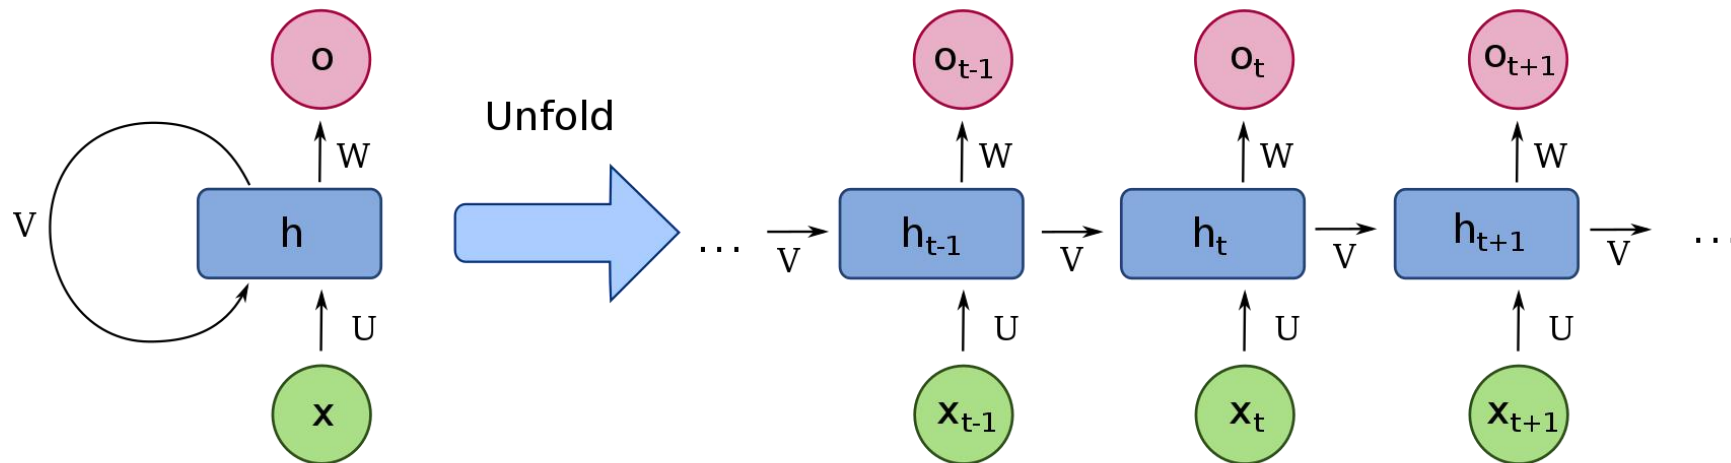

Inception V1



- <https://learnopencv.com/ultralytics-yolov8/>
- <https://jonascleveland.com/best-image-classification-models/>
- VGG-Net (Visual Geometry Group)
<https://medium.com/@siddheshb008/vgg-net-architecture-explained-71179310050f>

Deep Recurrent Neural Nets (RNN)

- Temporal sequence
- For example, for processing video or understanding natural language (e.g. English sentences)



[[Recurrent neural network – Wikipedia](#)]

Transfer Learning

- But... Training DL models to solve new problems is not easy. Training Neural Nets requires A LOT of labelled data.
 - Need big data and computing power
 - Access to lots of computing power (usually GPU processing)
- Good news!
 - “Extending a network’s capability to new kinds of images that weren’t part of the previous learning means transferring the knowledge to this new problem (transfer learning).”
 - “Use the majority of the layers of the network as they are (you freeze them) and then work on the final, output layers (**fine-tuning**)”
 - For example, use InceptionV3 feature extraction, but learn a new set of categories

Summary

- **Image segmentation** is the task of finding groups of pixels that “go together”.
- In salient **object segmentation**, the goal is to group pixels that belong to the same object
- Many methods exist for segmentation, each with its strengths and weaknesses. Recent methods based on Deep Learning have a good performance.

Summary (cont.)

- Computer Vision problems can be solved using three approaches: 1) Start from intuition and design a method, or 2) Start from data and use Machine Learning (ML) 3) Start from LOTS of data and use **Deep learning**, a new ML technique inspired by biological brains.
- Artificial Neural Networks solve AI / Machine Learning problems by building an artificial “brain” made of artificial neurons.
- An artificial neuron or **perceptron** can be trained to perform linearly separable binary classification. However, interesting problems require complex artificial neural networks often in a deep (multi-layer) architectures.

Summary (cont.)

- Until recently, training these networks was the challenge.
- Deep Learning methods offer improvements to **backpropagation** training and remove issues such as the **vanishing gradient** problem.
- Trained neural networks, such as **CNNs** for object detection in images, or **RNNs** for natural language processing, are often shared. These networks can be further trained or **fine-tuned** for customization.
- With the success of Deep Learning methods in training Neural Nets, many problems in ML and Computer Vision were solved!
- One specific architecture, called Convolutional Neural Net (CNN), implements convolution (or filtering) and thus feature detection, and is very effective in solving object detection, among other things.

References

- [1] Computer Vision: Algorithms and Applications, R. Szeliski
(<http://szeliski.org/Book>)
- [2] Learning OpenCV 3, A. Kaehler & G. Bradski
 - Available online through Safari Books, Seneca libraries
 - https://senecacollege-primo.hosted.exlibrisgroup.com/primo-explore/fulldisplay?docid=01SENC_ALMA5153244920003226&context=L&vid=01SENC&search_scope=default_scope&tab=default_tab&lang=en_US
- [3] Practical introduction to Computer Vision with OpenCV, Kenneth Dawson-Howe
 - Available through Seneca libraries
 - https://senecacollege-primo.hosted.exlibrisgroup.com/primo-explore/fulldisplay?docid=01SENC_ALMA5142810950003226&context=L&vid=01SENC&search_scope=default_scope&tab=default_tab&lang=en_US

References (cont.)

[4] Movahedi, V., “Automatic Extraction of Closed Contours Bounding Salient Objects: New Algorithms and Evaluation Methods”, Ph.D. Dissertation, York University, 2015.

[5] Howse, J., “Learning OpenCV 4 Computer Vision with Python 3”, Packt Publishing, 2020.

Readings

- Chapter 4 [2]
- Chapter 12 [3]

- [1] **A Practical Introduction to Computer Vision with OpenCV** by Kenneth Dawson-Howe Available online via library.senecacollege.ca: [Permalink](#)
- [2] **Learning OpenCV 4 Computer Vision with Python 3** by J. Howse & J. Minichino Available online via library.senecacollege.ca: [Permalink](#)
- [3] **Learning OpenCV 3** by A. Kaehler & G. Bradski Available online via library.senecacollege.ca: [Permalink](#) [4] **Mastering OpenCV with practical computer vision projects** by D.L. Baggio et al. Available online via library.senecacollege.ca: [Permalink](#)

Blobs

- A group of connected pixels in an image, sharing some property, e.g. grayscale or color

OpenCV SimpleBlobDetector

```
using namespace cv;
// Read image
Mat im = imread("blob.jpg", IMREAD_GRAYSCALE);

// Set up the detector with default parameters.
SimpleBlobDetector detector;

// Detect blobs.
std::vector<KeyPoint> keypoints;
detector.detect(im, keypoints);

// Draw detected blobs as red circles.
// DrawMatchesFlags::DRAW_RICH_KEYPOINTS flag ensures the size of the circle corresponds to the
// size of blob
Mat im_with_keypoints;
drawKeypoints(im, keypoints, im_with_keypoints, Scalar(0, 0, 255),
DrawMatchesFlags::DRAW_RICH_KEYPOINTS);

// Show blobs
imshow("keypoints", im_with_keypoints);
waitKey(0);
```

Older version n?

Use :

https://raw.githubusercontent.com/opencv/opencv/master/samples/cpp/detect_blob.cpp


<https://stackoverflow.com/questions/30622304/opencv-3-blob-detection-the-function-feature-is-not-implemented-in-detectand>

This class performs several filtrations of returned blobs. You should set `filterBy*` to true/false to turn on/off corresponding filtration. Available filtrations:

- **By color.** This filter compares the intensity of a binary image at the center of a blob to `blobColor`. If they differ, the blob is filtered out. Use `blobColor = 0` to extract dark blobs and `blobColor = 255` to extract light blobs.
- **By area.** Extracted blobs have an area between `minArea` (inclusive) and `maxArea` (exclusive).
- **By circularity.** Extracted blobs have circularity ($\frac{4*\pi*Area}{perimeter*perimeter}$) between `minCircularity` (inclusive) and `maxCircularity` (exclusive).
- **By ratio of the minimum inertia to maximum inertia.** Extracted blobs have this ratio between `minInertiaRatio` (inclusive) and `maxInertiaRatio` (exclusive).
- **By convexity.** Extracted blobs have convexity (area / area of blob convex hull) between `minConvexity` (inclusive) and `maxConvexity` (exclusive).

Default values of parameters are tuned to extract dark circular blobs.

[https://docs.opencv.org/3.0-last-rst/modules/features2d/doc/common_interfaces_of_feature_detectors.html#simpleblobdetector]



[<https://www.learnopencv.com/blob-detection-using-opencv-python-c/>]


Regional methods

Homogeneity Measure

- How do you measure if a region contains similar pixels?
 - Can use brightness, color, texture, etc.
 - Example: gray-level variance
1. Define a homogeneity measure (H)
 2. Set a threshold (t)
- If for a region R , $H(R) > t$, then pixels are similar, and belong to same segment
 - Otherwise, pixels do not belong to the same segment

Split-only Segmentation

- Start from the whole image as one region
- Loop over all (new) regions
 1. Calculate the **homogeneity** measure
 2. If the measure < threshold (did not pass the test), split the region (e.g. into 4 regions)
- Continue until all regions pass the homogeneity test



Example: Merge-only

- Start from the lowest level (e.g. each pixel being a region or each $N \times N$ neighborhood as a region)
- Loop over all (new) regions
 1. Try merging the region with neighbors
 2. Calculate the homogeneity measure
 3. Keep the merged region, if the measure $>$ threshold (passing the test)

Split & Merge Techniques[5]

- Procedure:
 1. Split the image into equally sized regions
 2. Calculate the homogeneity measure for each region
 3. If the measure $>$ threshold (passed the test), attempt to merge with its neighbor(s). If not, split the region
 4. Continue until all regions pass the homogeneity test

FloodFill flags

- Possible values for flag:

Low 8 bits	4	use 4-connectivity
	8	use 8-connectivity
High 8 bits	<code>cv::FLOODFILL_FIXED_RANGE</code>	A pixel will be compared to the original seed rather than to its neighbors
	<code>cv::FLOOD_MASK_ONLY</code>	The input image will not be modified, instead a mask argument will be the output (see [2] page 362)
Middle 8 bits	A numerical value	The value to use when filling the mask (default is filling with 1s)

- Use OR to set multiple flags:

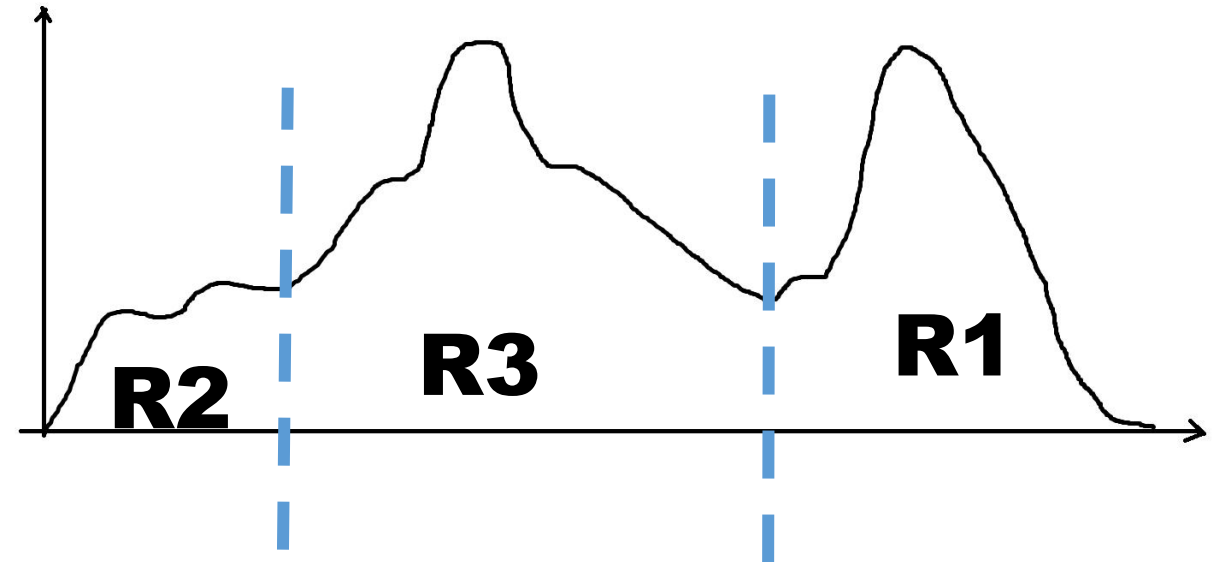

```
flags = 8  
| cv::FLOODFILL_MASK_ONLY  
| cv::FLOODFILL_FIXED_RANGE  
| (47<<8);
```

Clustering methods

- Similar ideas of splitting and merging, but not applied in the spatial (row and column) domain
- Instead other domains, such as the histogram
- Therefore can group pixels that are not connected
- Suitable for **occlusion** (when an object is partially covered)

Histogram-based methods

- Simple example:
 - Start as the entire image as one region
 - Loop until reaching the expected number of regions
 - Find a peak and put thresholds on either side of the peak
 - Use the thresholds to split the region into two regions




K-means clustering

- A method for clustering (grouping) data into k groups
- Needs a measure of similarity/ distance between two data points
 - Examples:
 - Color similarity between two pixels
 - Spatial distance between two pixels
 - Or more complex measures
- Procedure
 - Initialize k centers or means (m_1, m_2, \dots, m_k)
 - Loop until no change
 - *Assignment step*: Assign all data to the most similar/ nearest center
 - *Update step*: Calculate the new center/mean as the mean of data belonging to the same cluster



K-Means Clustering...




**...clearly
explained!!!**

<https://youtu.be/4b5d3muPQmA>

Example- 2 dimensional data

- [<http://dni-institute.in/blogs/k-means-clustering-algorithm-explained/>]



Other segmentation methods

Graph methods

- Regions are modelled as the nodes in a graph.
- Neighboring regions are connected by edges in the graph
- The weights of the edges is defined by the similarity between the regions
- Using graph optimization methods to group the nodes
- For example:
 - Graph-cut
 - Min-cut
 - Grab-cut

Segmentation Algorithms [4]

- Regional methods
 - Optimize the labelling of pixels
- Active Contour methods
 - Find the optimal deformation and location of a deformable contour around an object in the image
- Contour grouping methods
 - Search for the optimal grouping of boundary entities (e.g. edgels, line segments or curvelets) to form an object contour
- Deep Learning methods
 - Using Neural Networks and training them using lots of data

OpenCV findContours

findContours

Finds contours in a binary image.

```
C++: void findContours(InputOutputArray image, OutputArrayOfArrays  
contours, OutputArray hierarchy, int mode, int method, Point  
offset=Point())
```

```
C++: void findContours(InputOutputArray image, OutputArrayOfArrays  
contours, int mode, int method, Point offset=Point())
```

[https://docs.opencv.org/3.0-last-rst/modules/imgproc/doc/structural_analysis_and_shape_descriptors.html]

```
namedWindow( "image", 1 );
imshow( "image", img );
//Extract the contours so that
vector<vector<Point> > contours0;
findContours( img, contours0, hierarchy, RETR_TREE, CHAIN_APPROX_SIMPLE);

contours.resize(contours0.size());
for( size_t k = 0; k < contours0.size(); k++ )
    approxPolyDP(Mat(contours0[k]), contours[k], 3, true);
```

```
drawContours( cnt_img, contours, _levels <= 0 ? 3 : -1, Scalar(128,255,255),
              3, LINE_AA, hierarchy, std::abs(_levels) );

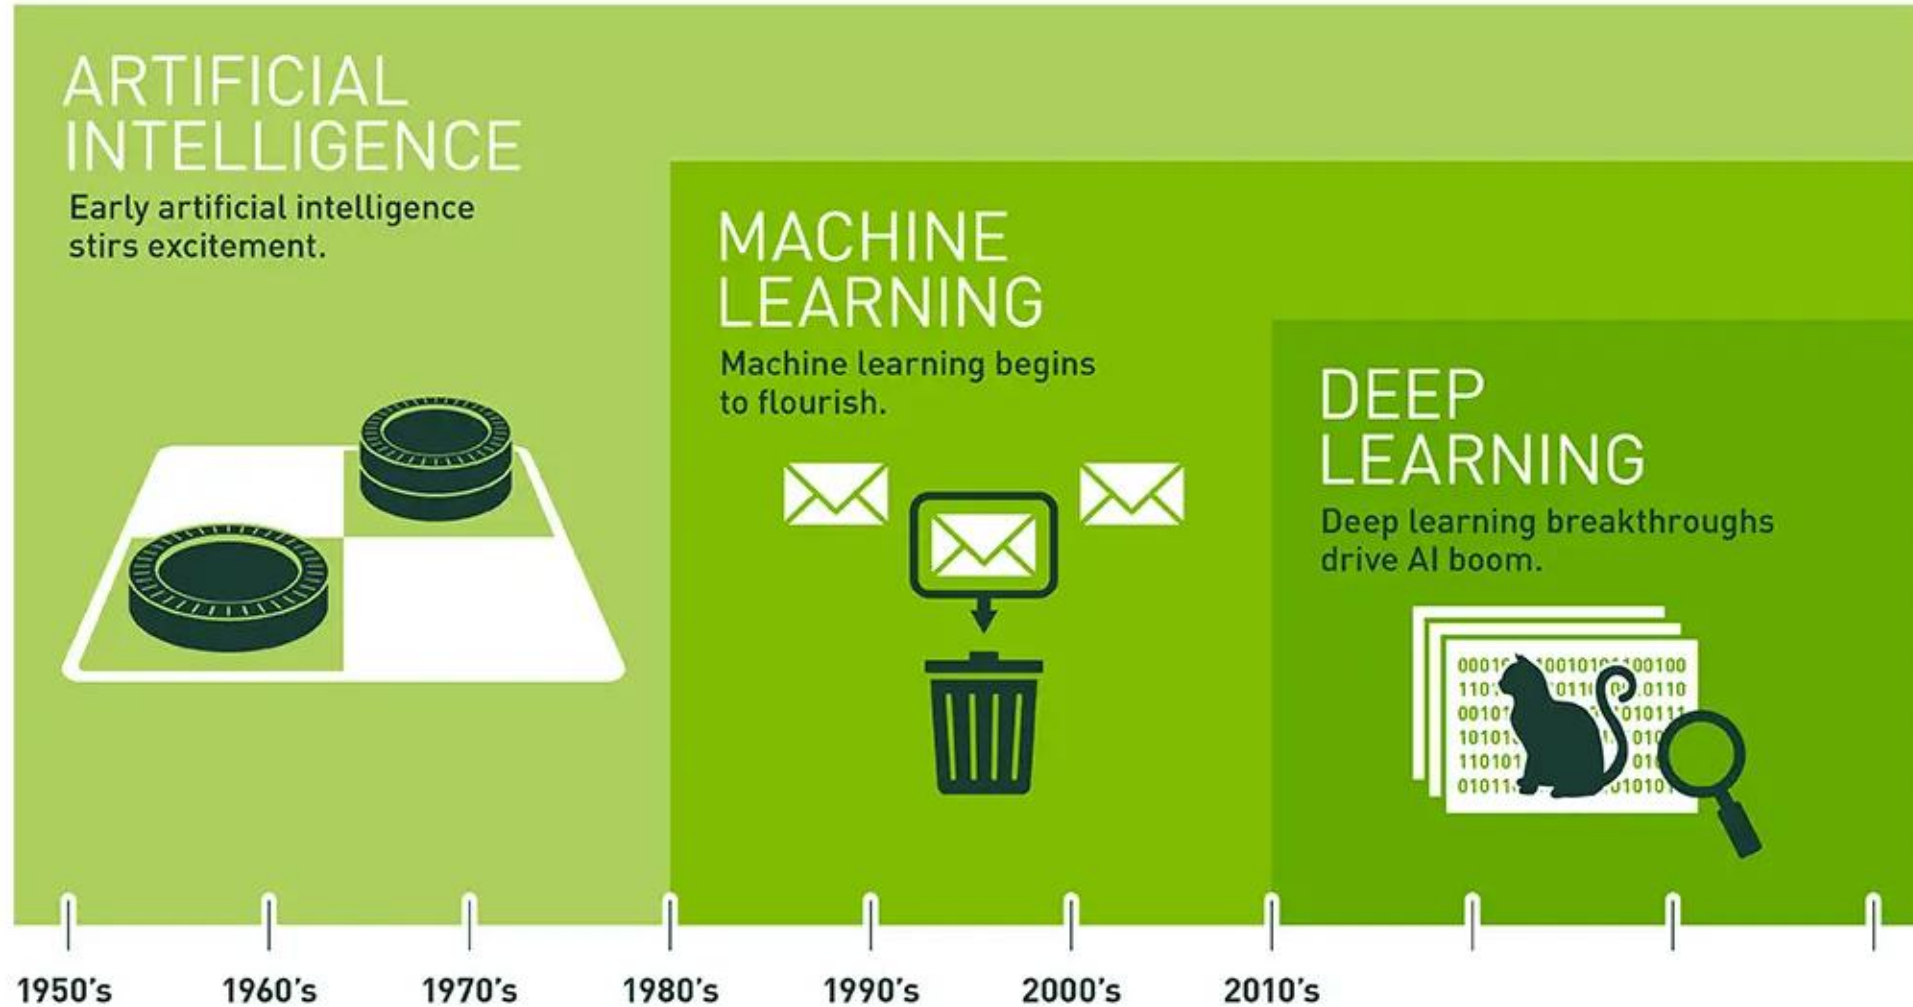
imshow("contours", cnt_img);
```

[<https://github.com/opencv/opencv/blob/master/samples/cpp/contours2.cpp>]

Iterative training & epochs


- Backpropagation
 - Mathematical formulation for using feedback from trial and error to train the network
- **Epochs:** Repeated iterations of seeing training samples
- Learn from mistakes!
- The **error** calculated at the last layer is sent back through the network and used for calculating adjusted weights

Timeline



[<https://blogs.nvidia.com/blog/2016/07/29/whats-difference-artificial-intelligence-machine-learning-deep-learning-ai/>]

Artificial Neuron



Transfer function: $y_i = w_0 + w_1x_1 + \dots + w_nx_n$

Activation function: $z_i = \frac{1}{1+e^{-y_i}}$
the logistic or sigmoid function (above)
Or other squashing functions

Single Neuron is a linearly weighted sum, followed by a “squashing” function. z_1 has range $[0,1]$

Computing for Deep Learning

- Huge amounts of data needed for training deep nets
- Many iterations (epochs) needed
- Not feasible using normal computers
- Improvements in Computing
 - GPU computing for images
 - Cloud/clusters (Google, Facebook, Microsoft, IBM)
- Example
 - Google Brain project, 16,000 computers, more than a billion weights, for unsupervised learning from YouTube videos

Examples of DL models & applications

- Image classification, object detection
- Visual Object Tracking
- Automatic Speech Recognition
- Natural Language Processing
- Recommendation
 - Netflix
 - Amazon

Limitations of DL

- Needs lots and lots of data
- If supervised learning, all need to be labelled
- Needs lots of computing power/ time
- DL models do not generalize
 - If only seen samples of cats and dogs, they won't (necessarily) generalize to learn a class of animals
- Not easy to explain their output (how and why)

Examples

- Deep Learning In 5 Minutes: <https://youtu.be/6M5VXKLf4D4>
- Top Deep Learning Projects: <https://youtu.be/N082bM72byY>
- 9 Cool Deep Learning Applications: <https://youtu.be/Bui3DWs02h4>
- 10 More Cool Deep Learning Applications:
<https://youtu.be/hPKJBXkyTKM>